

# Data deep-dive research

Seungwoo Lim

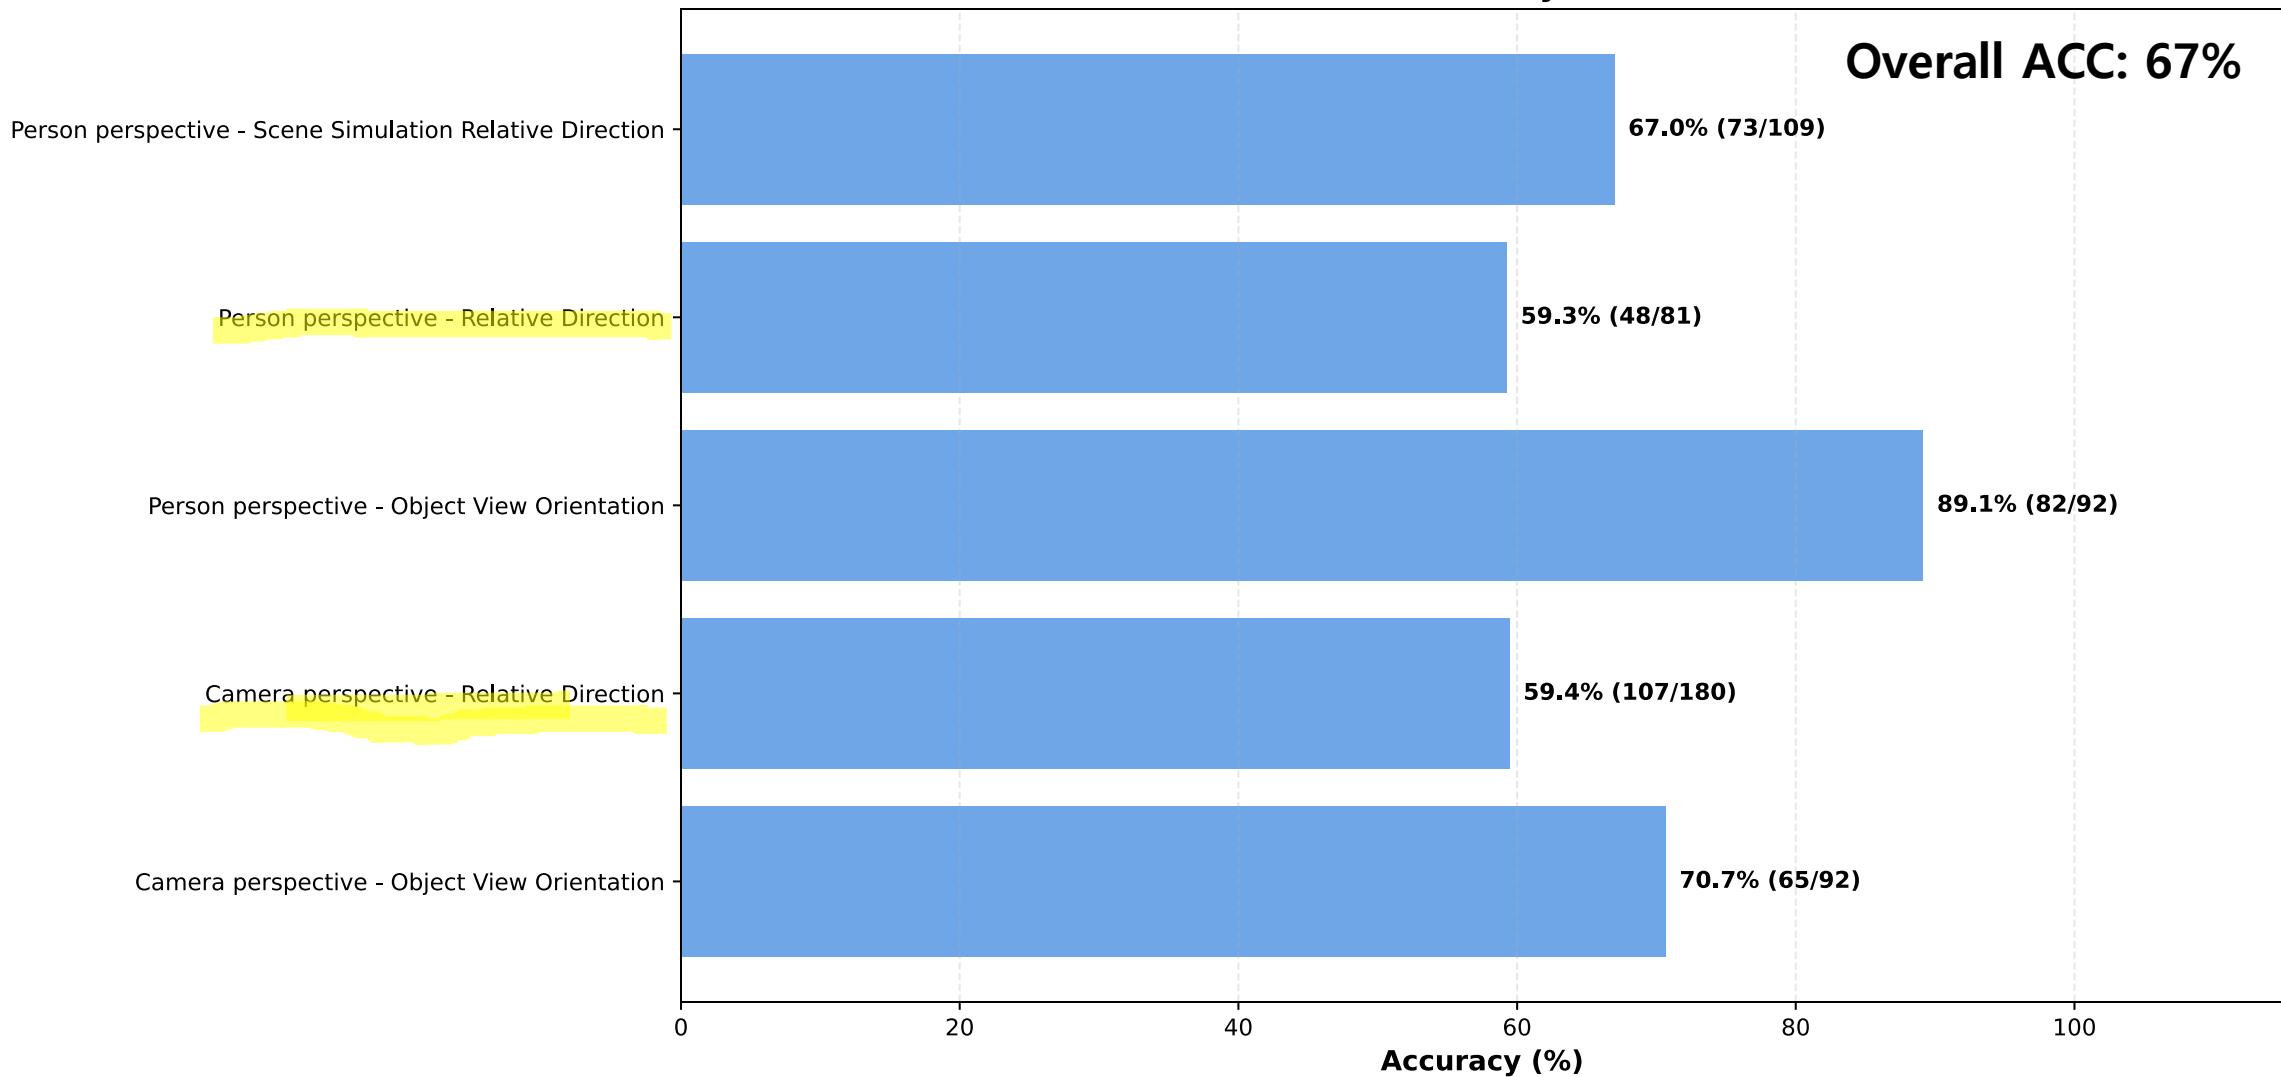
# Short recap

- Problems with the last research
  - Haven't looked into the data in detail.
  - A one-to-one comparison of a different model(MVSM) from the paper

# Introduction

- Is there a task that you get more wrong than others? Why?
- If you get VQA wrong, is there a tendency for the model to give incorrect answers?
  - Left -> Right? Or is it completely wrong?
- Baseline : Vanilla MVSM (fine tuned Qwen2.5)

### Task-wise Accuracy (Baseline)



# Hypothesis 1

## Data Augmentation

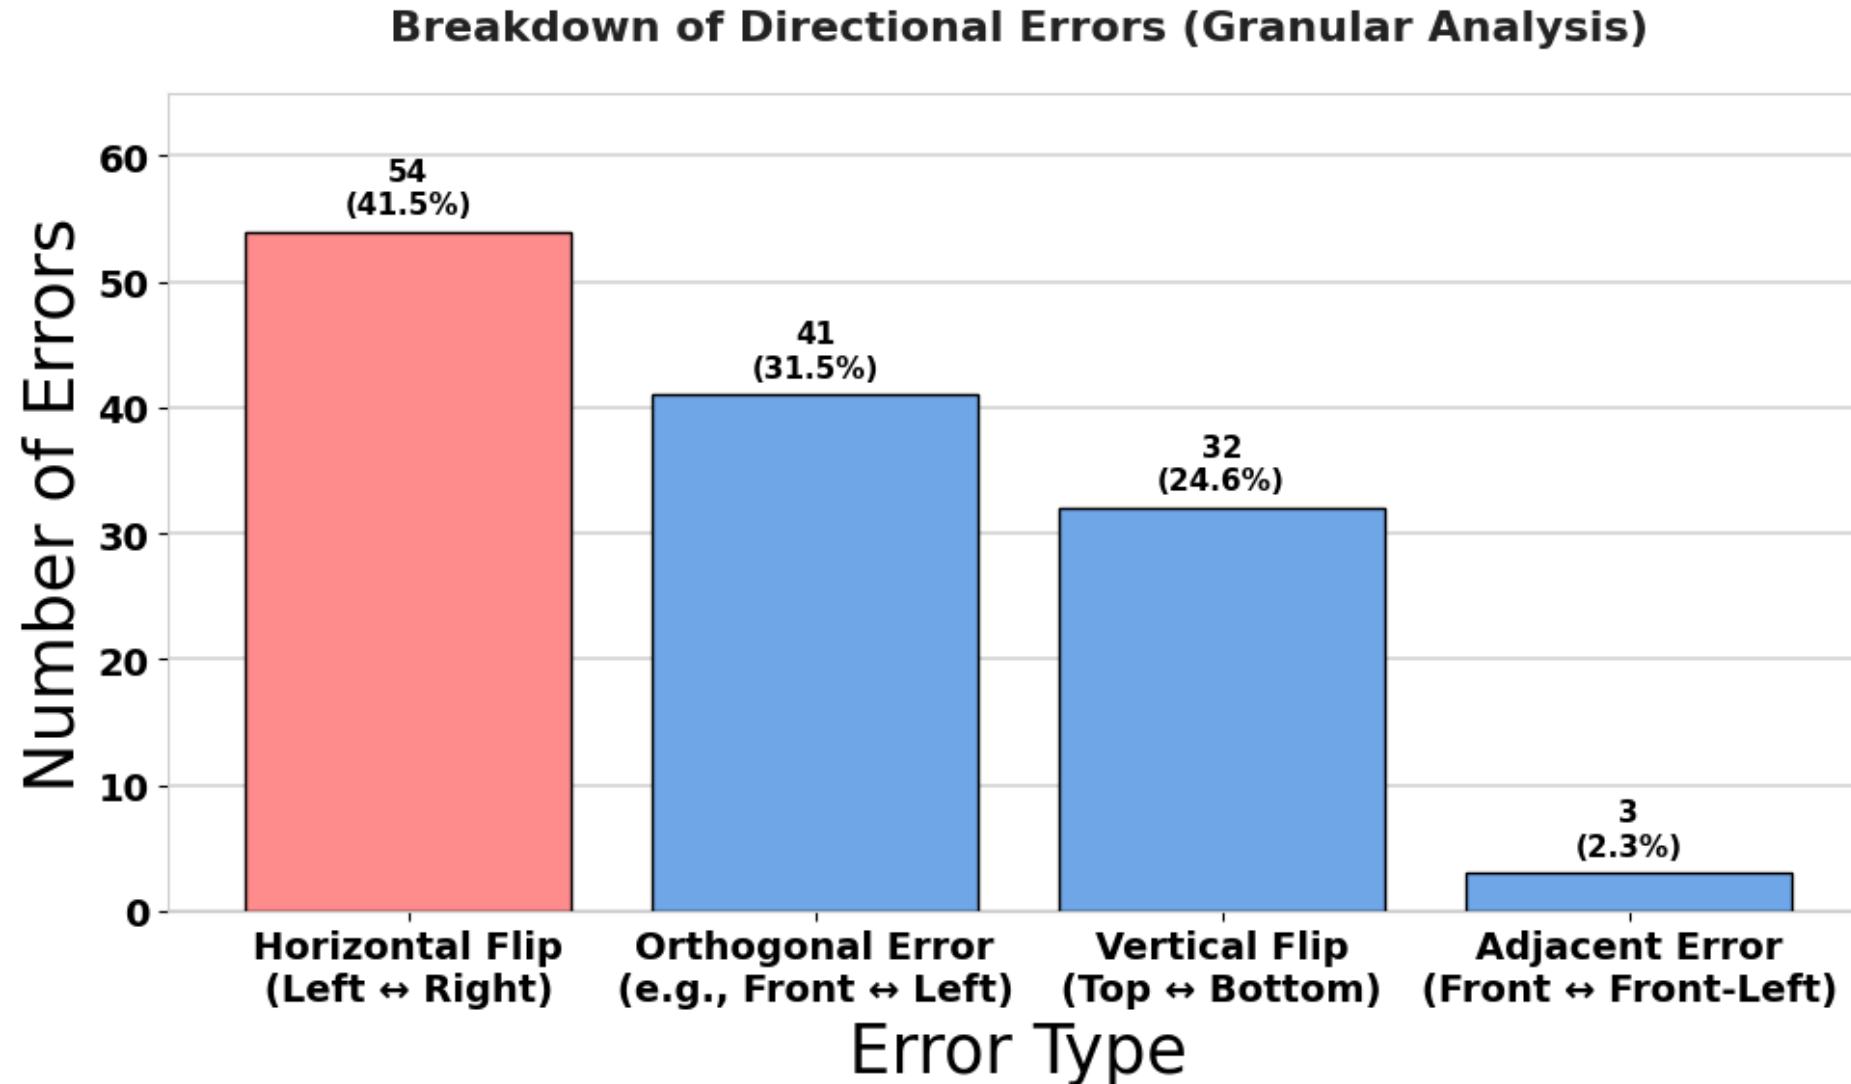
# Hypothesis 1 : Data augmentation

- As demonstrated in Vision-Language research, focusing on **learning hard negatives** that confuse the model is key to improving performance, rather than simply augmenting data.
- Hard negative :
  - The most common incorrect answer the model gets wrong
  - image-text pair that is semantically similar but differs in fine-grained details.

# Hypothesis 1 : Data augmentation

- **Horizontal Errors:**
  - (Front/Back)Left  $\leftrightarrow$  (Front/Back)Right
- Orthogonal Errors:
  - Front  $\leftrightarrow$  Left/Right
- Vertical Flip:
  - Top  $\leftrightarrow$  Bottom
- Adjacent Error:
  - Angles that are not 90 or 180 degrees
  - Front  $\leftrightarrow$  Front-left

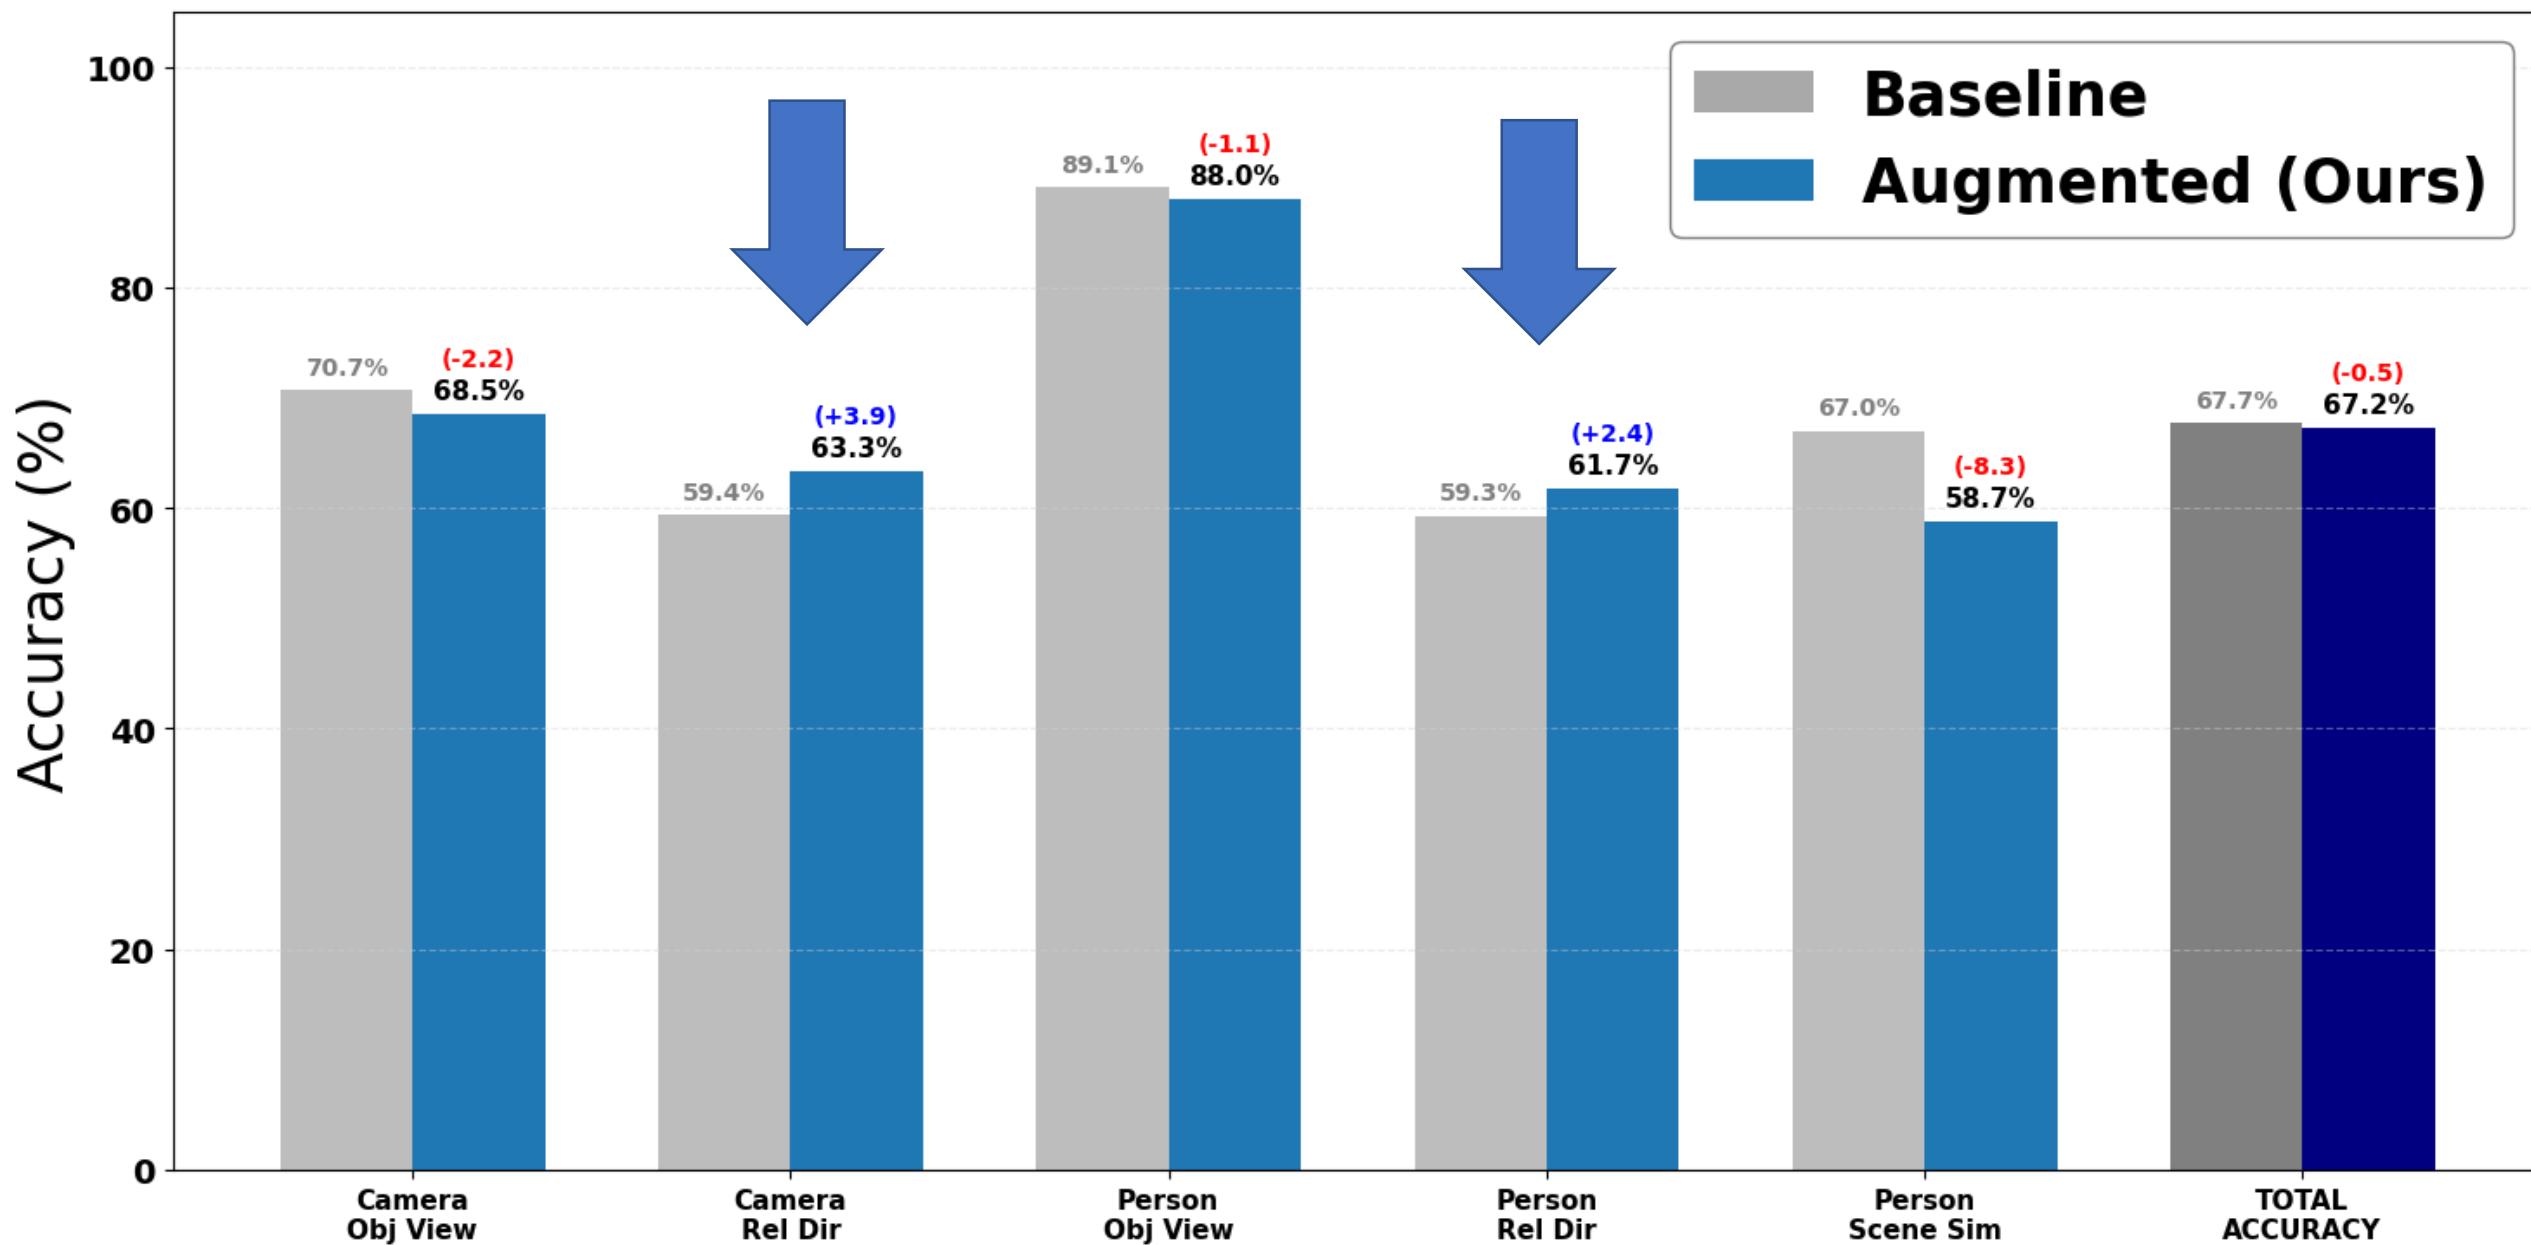
# Problem 1 : Direction analysis



# Hypothesis 1

- Augment Train dataset
  - 4548(baseline train data) -> 9048 (augmented train data)
  - Flip the image, flip the label( left , right)

# Benchmark Performance Comparison (Baseline vs Augmented)



# Result 1 : Data Augmentation

- Although we achieved accuracy improvement in the targeted relative direction,
- We confirmed that the performance of other tasks actually decreased.
  - Which is called **Catastrophic Forgetting\***
  - -> To solve this, we can try “cognitive replay”
  - It is a method of repeatedly learning previous data.

# Hypothesis 2

Bounding box + Visual prompting

# Problem 2 : Object detection fail



☒ **WRONG**

**[Question]**

**From the perspective of the boy, where is the TV located?**

**[Options]**

- A. back-left B. front C. right D. left

# Problem 2 : Object detection fail

- The question is about "person perspective" but the model is likely looking at "the entire image"

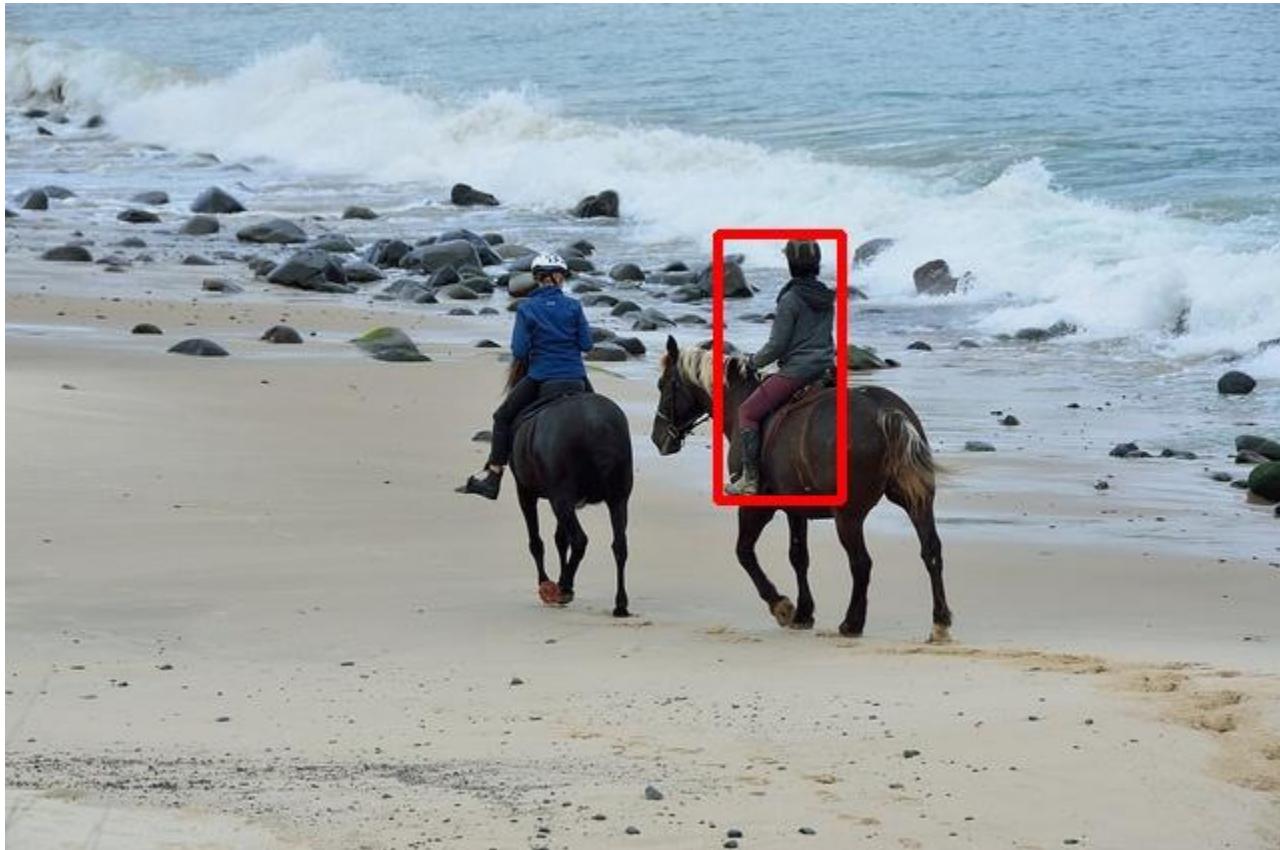
# Hypothesis 2 : Bounding box + Visual prompting

- The model may not be good at recognizing **reference objects**.
  - "From the perspective of **the boy**, where is the TV located?"
- -> Then let's make the reference point recognition more certain!
  - Maybe Bounding box might be helpful!
  - Used "Ultralytics YOLO-World"

# How does “Yolo world” works?

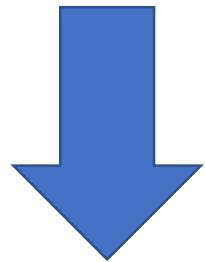
- Existing YOLO (v5, v8, etc.) could only find fixed, pre-trained classes (e.g. 80 classes including people, cars, dogs, cats, etc.).
- 1. Text Encoder : Convert user-entered text (e.g., "cat") into vector that the computer can understand. This process utilizes CLIP, a large-scale language-image model, to extract the meaning of the text.
- 2. Image Encoder : Analyzes the input image to extract visual features
- 3. Fusion : If a specific area of the image has a high similarity to the text vector, we determine that it is the object the user is looking for and draw a bounding box there.

# Example



*MS-CoCo  
Example*

From the perspective of the person wearing the grey clothes, where is the person wearing the blue clothes?



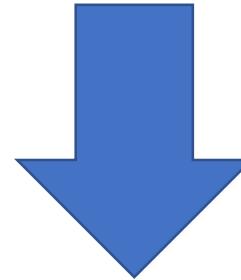
From the perspective of the person **in the red bounding box** wearing the grey clothes, where is the person wearing the blue clothes?

# Example



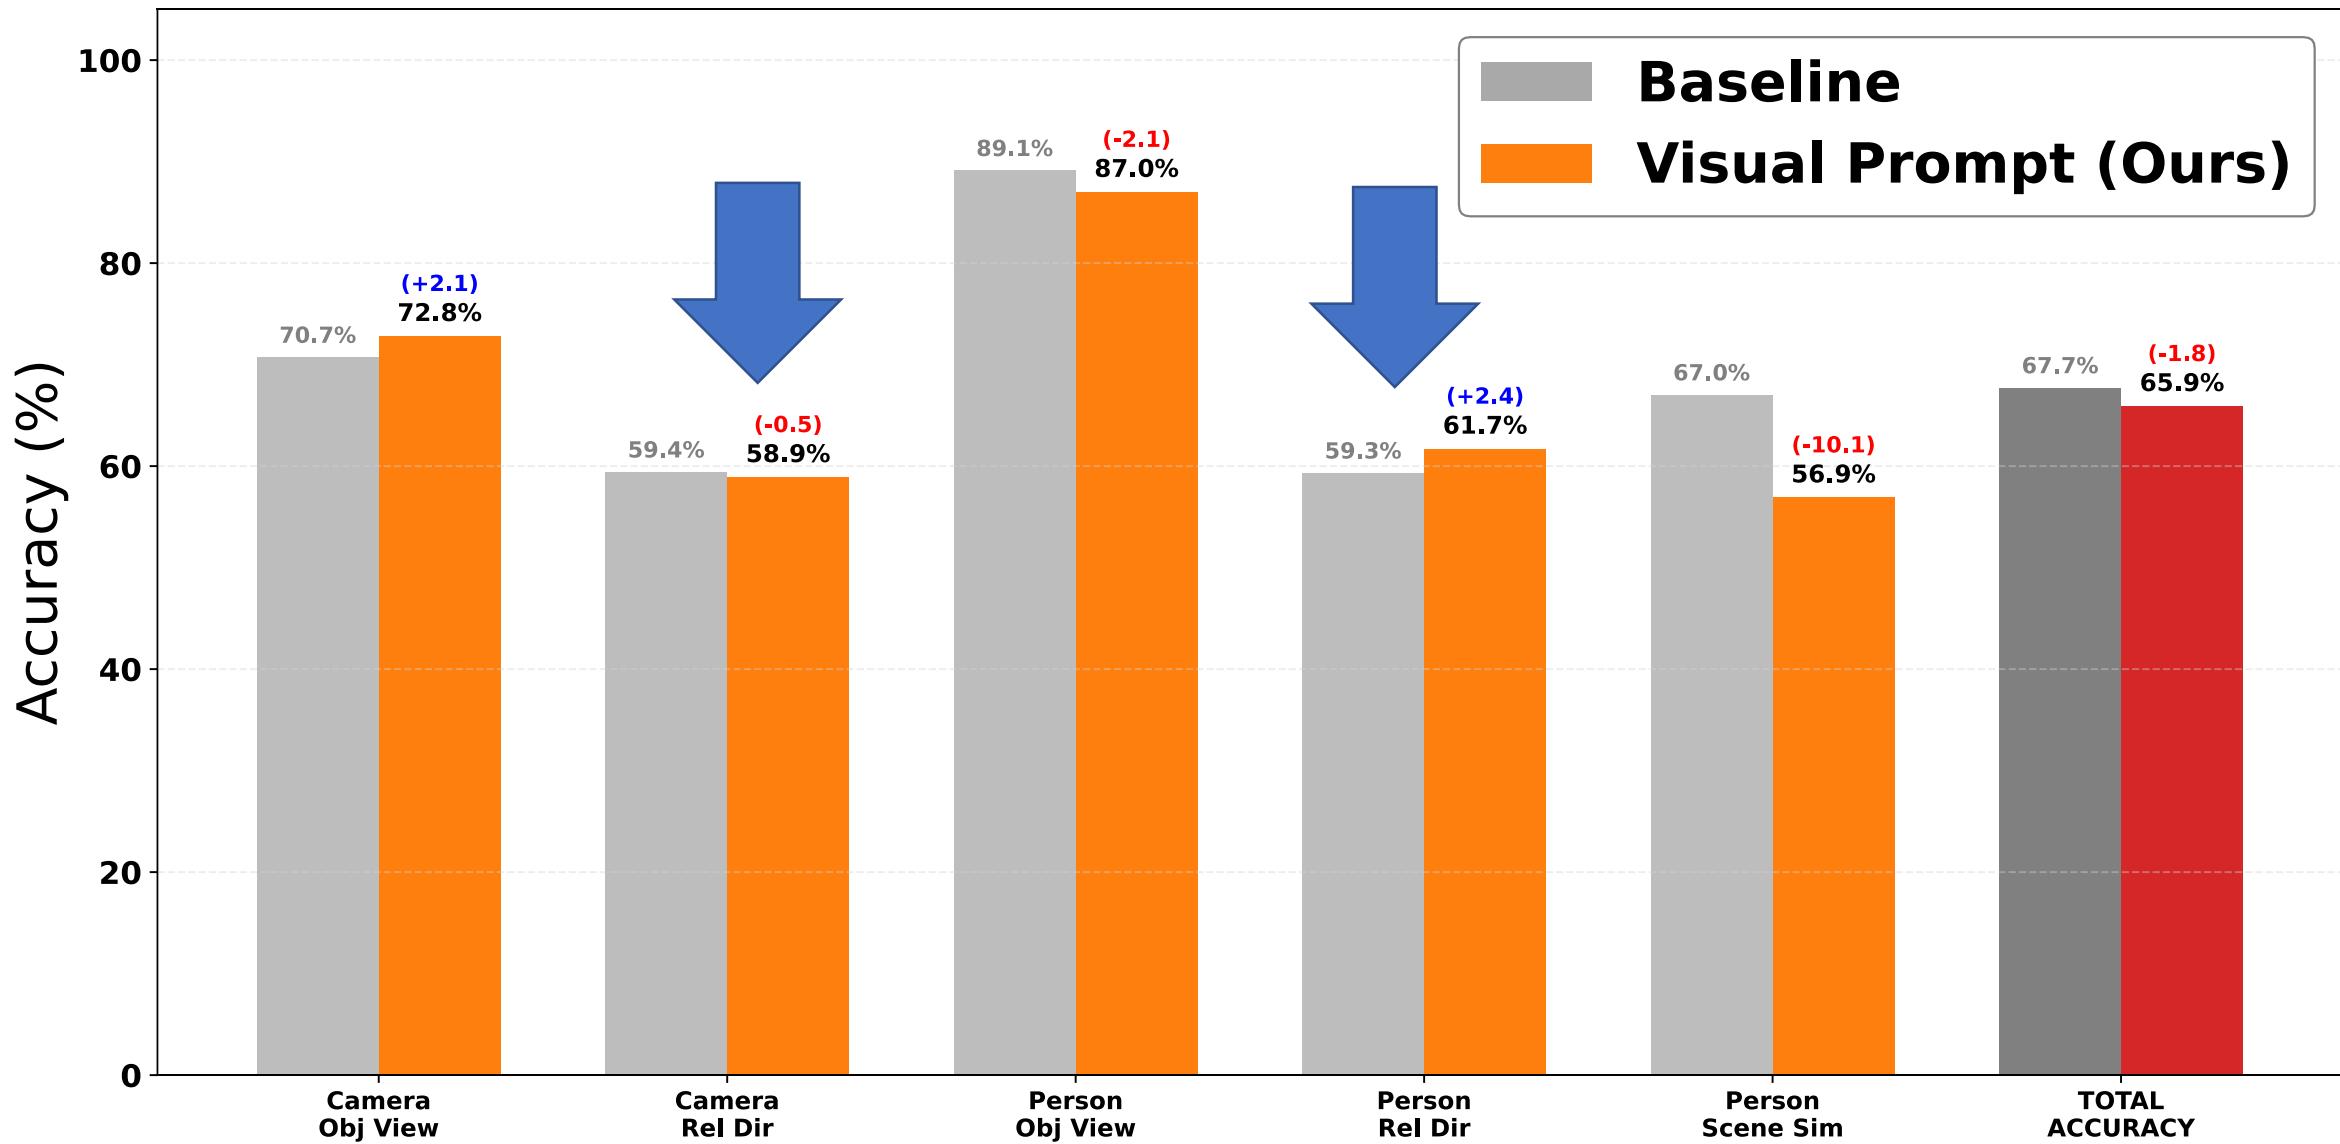
*ScanNet Example*

How is the sofa positioned with respect to the table?



How is the sofa positioned with respect to the table in the **red bounding box**?

# Performance Comparison: Baseline vs Visual Prompting



# Result 2 : Bounding Box

- Result :
  - [Per] Rel dir (+2.4%), [Camera] Obj + 0.1% : Performance improved in problems where the anchor point must be clear.
- Limitations:
  - [Per] Scene simulation (0.1%) : The surrounding environment (walls, obstacles, spatial structure) must be grasped as a whole, but this is interpreted as a kind of **tunnel vision** phenomenon where the model's gaze is trapped within a red box.

# Do bounding boxes induce tunnel vision?

- Used Saliency map

$$S_i = \left| \frac{\partial f(x)}{\partial x_i} \right|$$

$f(x)$  : Logit of  
 $x_i$  prediction  
 $x_i$  : pixel value

- A saliency map is a way to visualize how sensitive each element of the input is to the prediction by calculating the gradient of the input  $x$  with respect to the model output  $y$ .
- "If this pixel value changes even a little, how much does the correct answer (Logits) change?"
- Why not attention map?
  - It shows where the model "routed" during its computation, doesn't guarantee that it influenced the correct answer.

**[Rel Dir] From the perspective of the man who is looking at the computer, where is the man wearing a hat positioned?**



**Original Image**



**Original Saliency**



**Visual Prompt Image**



**Boxed Saliency**

**[Rel Dir] From the perspective of the woman wearing the green helmet,  
where is the person in black clothes?**



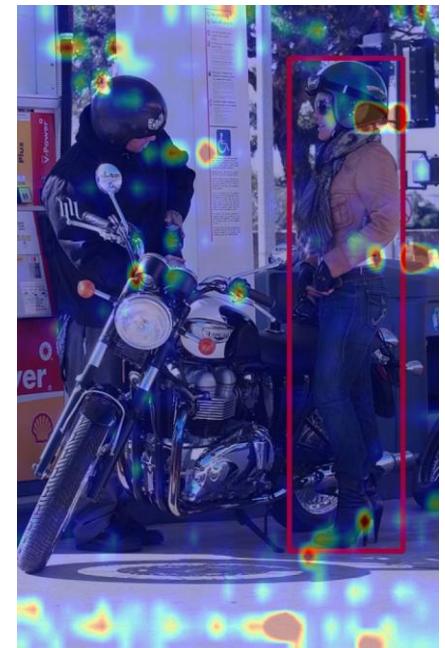
**Original Image**



**Original Saliency**



**Visual Prompt Image**



**Boxed Saliency**

# Hypothesis 3

Chain of thought

# Hypothesis 3 : Chain of Thought



Q : With the camera's viewpoint as the front, which direction is the elephant facing in the image? A. right B. front C. back-left D. left Answer with the option letter.

"A"



Let's think step by step to determine the spatial relationship.

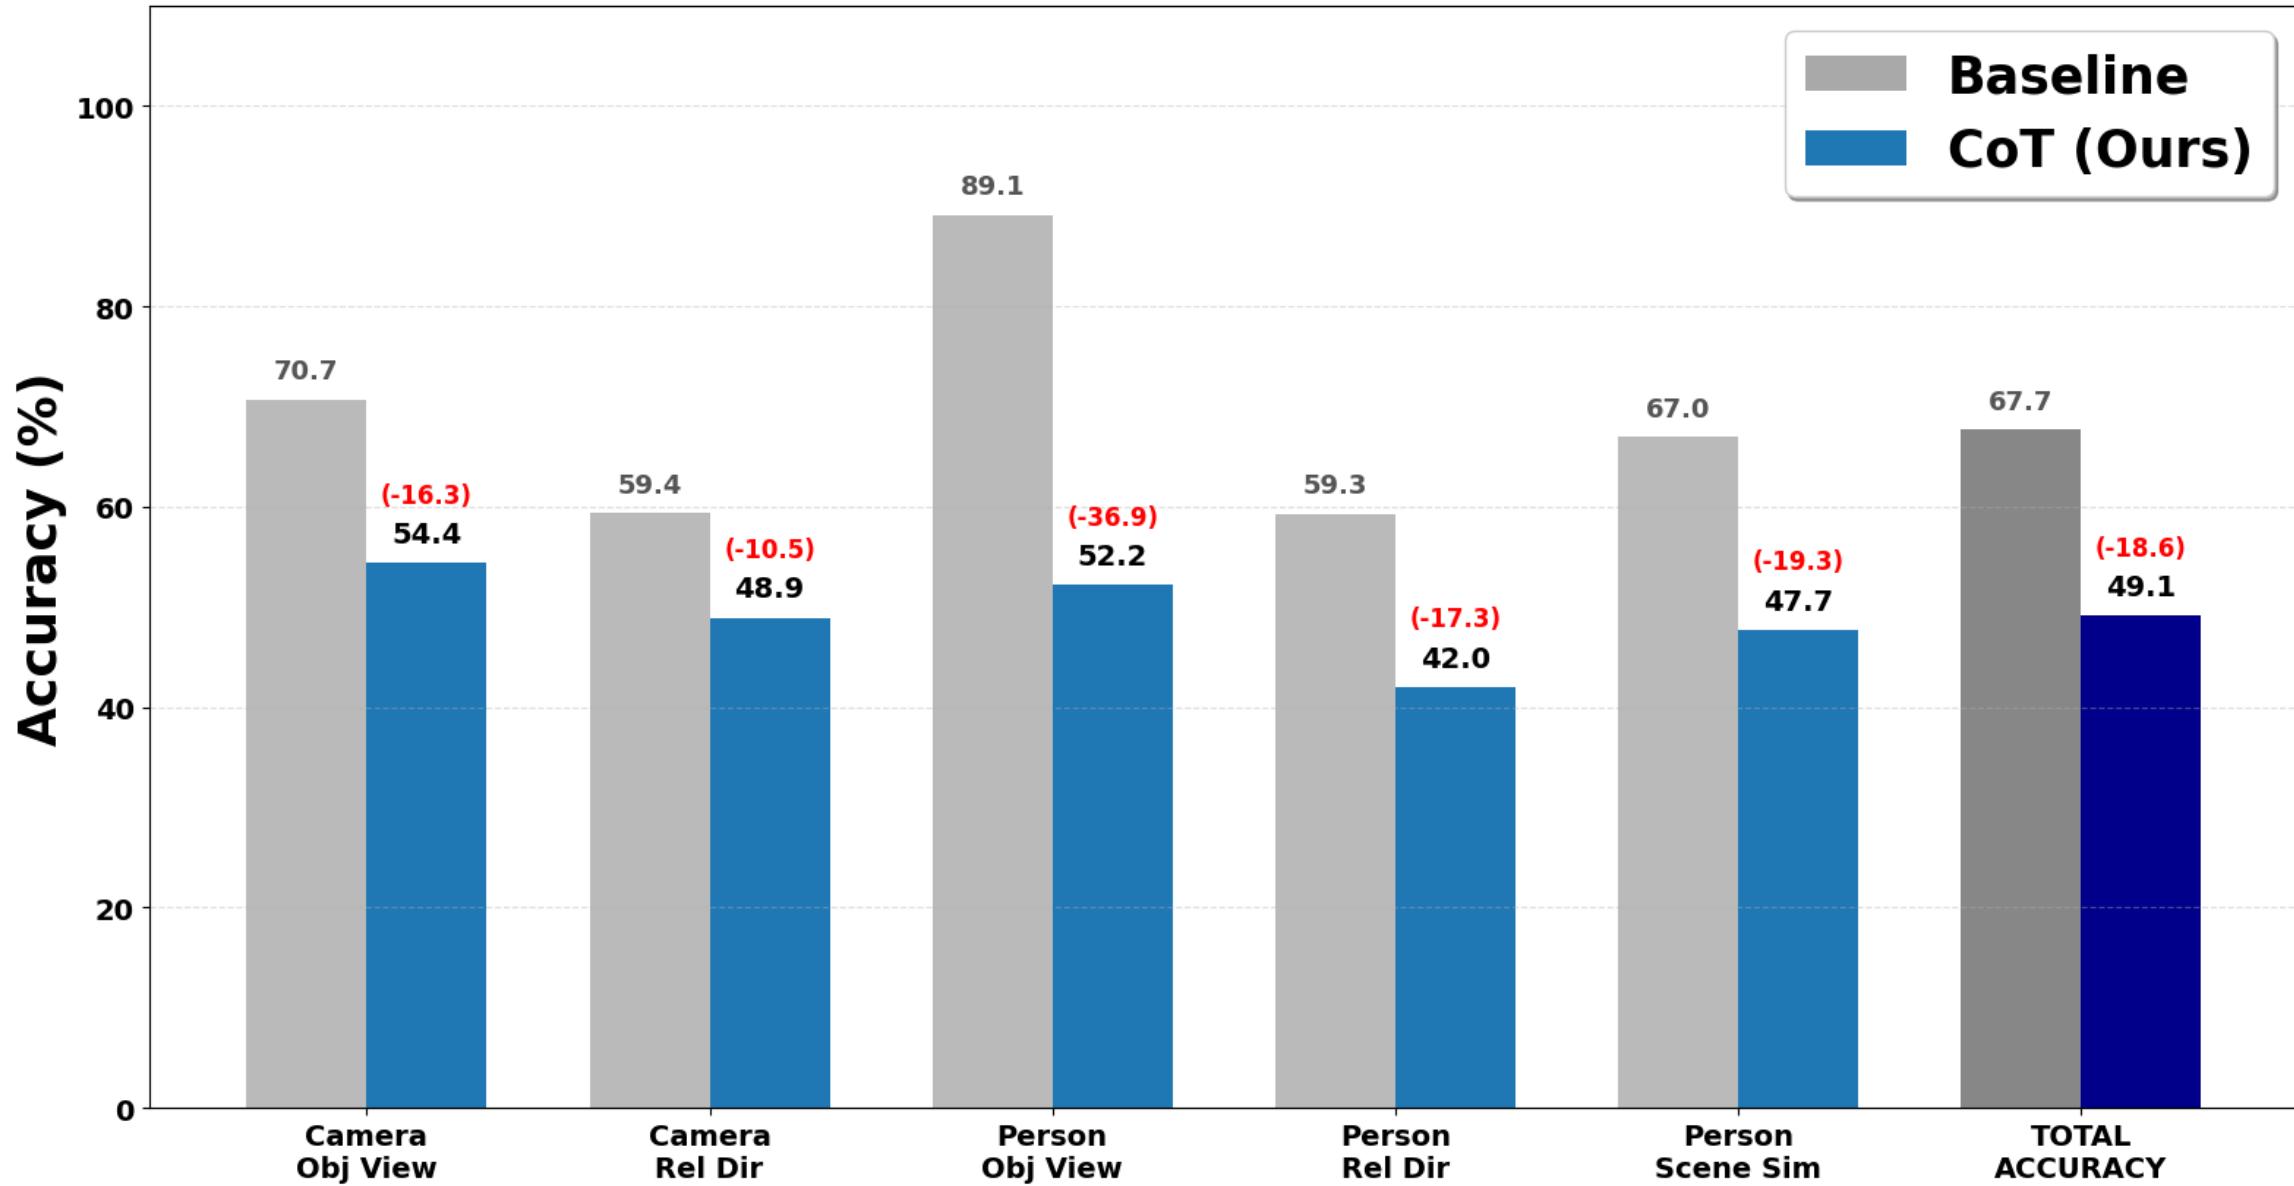
1. First, I identify the reference point: the camera viewpoint.
2. Next, I locate the target object: man.
3. By analyzing their relative positions in the 3D space, man is positioned to the right of the camera viewpoint.

Therefore, the correct option is **A**.

# Distilling the knowledge

- Used Qwen 2.5 7B model for **teacher** , and our Qwen 2.5 3B model will be a student
  - Qwen2.5 73B model was unusable due to OOM issue.

## Benchmark Performance: Baseline vs CoT(w Hard target)



# Result 3 : CoT

- Why is it so bad?
  - Used the Qwen 2.5 **7B** model instead of the 72B model or other higher end models.
  - Noise was mixed in during the process of creating **unnecessary forced logic**.
- How can we overcome this?
  - Better open source model (InternVL2-26B, Llama-3.2-11B-Vision, Qwen3)

# Conclusion

- We confirmed that data augmentation(for hard negatives) and bounding boxes helped improve performance, but the CoT experiment failed.

# Thank you!

daniel73919@gmail.com